

Computer System Organization (CS2.201)

Lecture # 05-07

Instruction Set Architecture / Assembly Language Programming


Avinash Sharma

Center for Visual Information Technology (CVIT),
IIIT Hyderabad

Slide content acknowledgment: Dr. Suresh Purini & other public sources


Intel Microprocessor Design History

- 8086: One of the first single-chip, 16-bit microprocessors.
- The 8088, a variant of the 8086 with a 20 bit address space, formed the heart of the original IBM personal computers.
- Architecturally, the machines were limited to a 655,360-byte address space as and the operating system reserved 393,216 bytes for its own use.



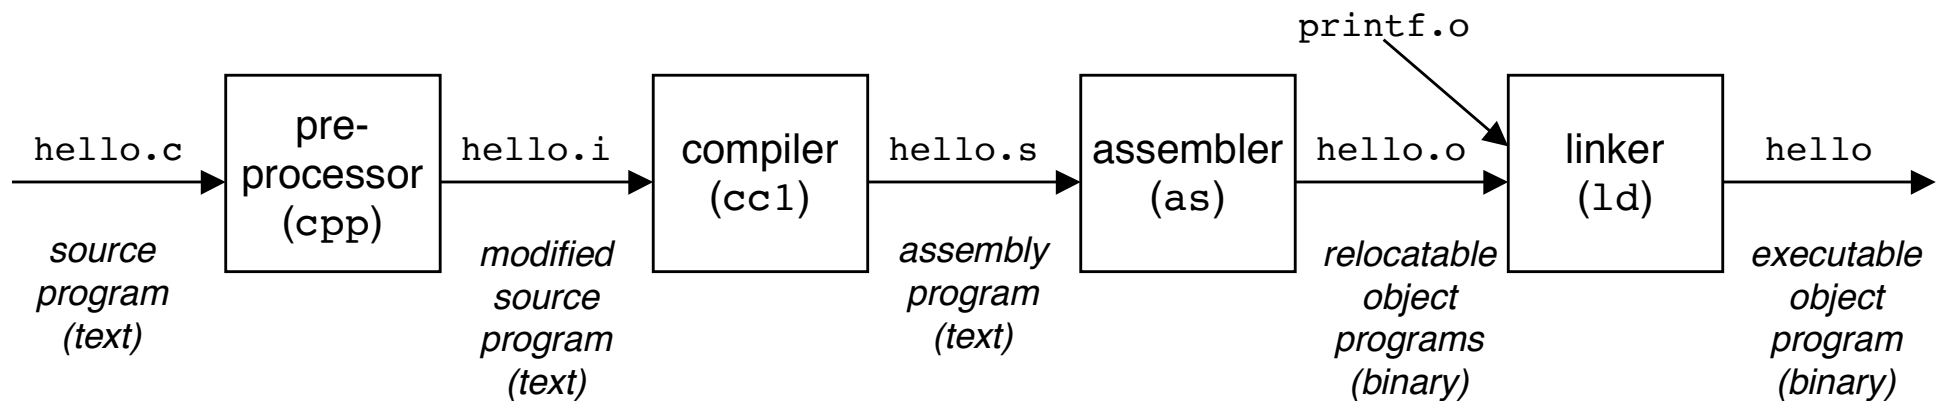
Intel Microprocessor Design History

- **80286:** More transistors.
- **i386:** Expanded to 32bits architecture. Added the flat addressing model used by Linux.
- **i486:** Integrated the floating-point unit onto the processor chip
- **Pentium:** Improved performance.
- **PentiumPro:** Introduced a radically new processor design, internally known as the *P6* microarchitecture. Added a class of “conditional move” instructions to the instruction set.



Typical Compilation System

```
gcc -o hello hello.c
```

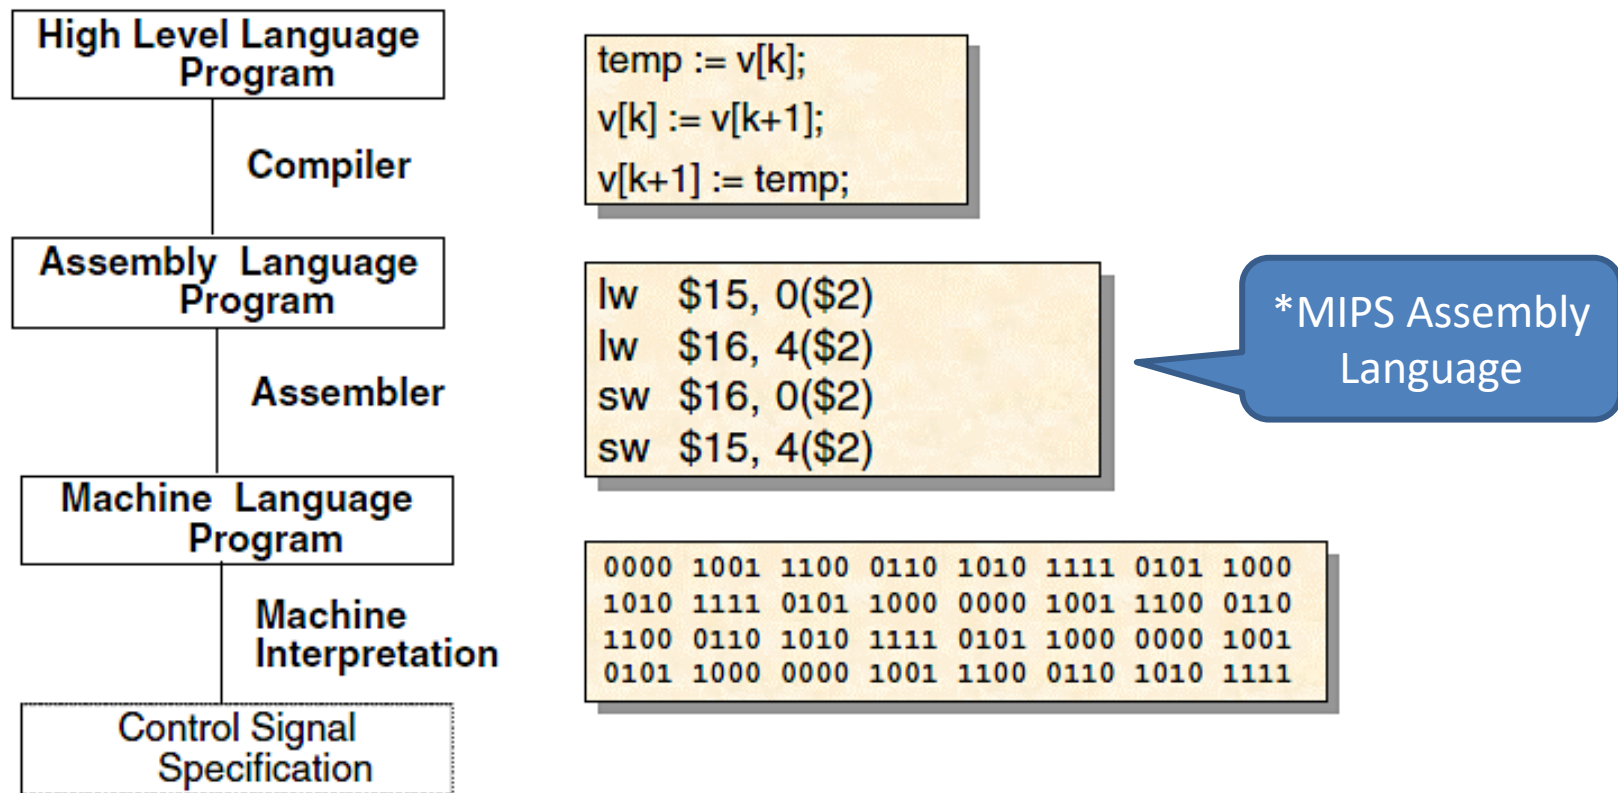


Programming Abstractions

We can program a microprocessor using

- a) Instruction opcodes (also called Machine Code)
 - b) Assembly language
 - c) High level programming languages
-
- ☐ The level of **abstraction** increases from Top to Bottom.
 - ☐ As the level of abstraction increases, ease of programmability also increases!
 - ☐ Hmm, but we may lose the fine-grained control over the underlying hardware?
-

Levels of Abstraction



*Microprocessor without Interlocked Pipeline Stages

Source: Prof. Cheung (ICL)

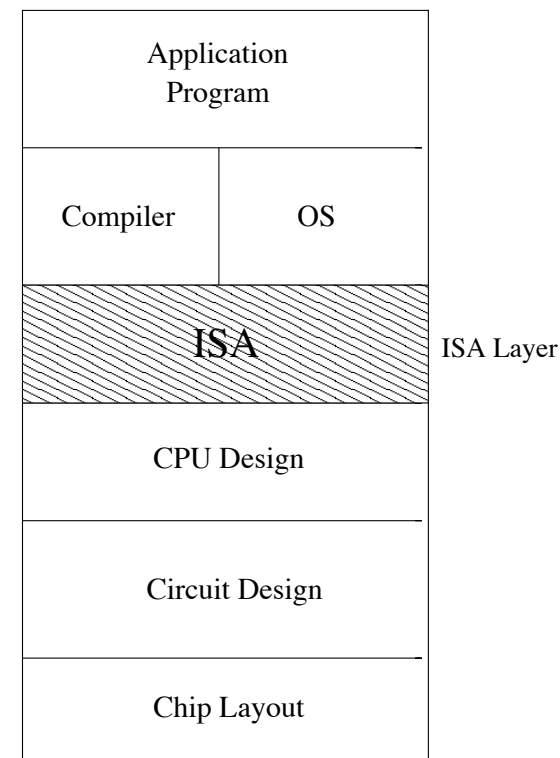
Instruction Set Architecture

Assembly Language View

- Processor state: registers, memory, etc.
- Instructions and how instructions are encoded

Layer of Abstraction

- Above: how to program machine, processor executes instructions sequentially
- Below: What needs to be built
 - Use variety of tricks to make it run faster, e.g., execute multiple instructions simultaneously
 - Safeguard to ensure that the overall behavior matches the sequential operation dictated by the ISA.



Instruction Set Architecture


- x86-64, describe the behavior of a program as if each instruction is executed in sequence, with one instruction completing before the next one begins.
 - The processor hardware is far more elaborate, executing many instructions concurrently, but they employ safeguards to ensure that the overall behavior matches the sequential operation dictated by the ISA.
 - Machine-level program use a virtual address space, providing a memory model that appears to be a very large byte array.
 - The actual implementation of the memory system involves a combination of multiple hardware memories and operating system software
-

Instruction Set Architecture

- There are various means of giving a semantics or meaning to a programming system.
 - Probably the most sensible for an assembly (or machine) language is an operational semantics, also known as an *interpreter semantics*.
 - That is, we explain the semantics of each possible operation in the language by explaining the effect that execution of the operation has on the machine state.
-

Instruction Set Architecture

- The most fundamental abstraction for the machine semantics for the x86 is the ***fetch-decode-execute*** cycle.
- The machine repeats the following steps :
 1. fetch the next instruction from memory (the PC tells you which is next);
 2. decode the instruction (in the control unit);
 3. execute the instruction, updating the state appropriately;
 4. go to step 1.



This is also called the von Neumann architecture.

Assembly Code Representation

- Very close to machine code.
 - More readable textual format, as compared to the binary format of machine code.
 - Understand assembly code and how it relates to the original C code is a key step in understanding how computers execute programs.
 - Enable visibility to key process states like *registers*, *program counter*, etc.
-

Instruction Format

- IA32 instructions can range in length from 1 to 15 bytes.
 - The instruction encoding is designed so that commonly used instructions and those with fewer operands require a smaller number of bytes as compare to less common ones or ones with more operands.
 - The instruction format is designed in such a way that from a given starting position, there is a unique decoding of the bytes into machine instructions.
-

Assembly Code Example

C Code (code.c)

```

1  int accum = 0;
2
3  int sum(int x, int y)
4  {
5      int t = x + y;
6      accum += t;
7      return t;
8  }

```

Parts of Assembler Output

```

sum:
    pushl    %ebp
    movl     %esp, %ebp
    movl     12(%ebp), %eax
    addl     8(%ebp), %eax
    addl     %eax, accum
    popl     %ebp
    ret

```

gcc -O1 -S code.c

Disassembled Output

```

1  00000000 <sum>:
    push    %ebp
2  0: 55 89 e5 0c 8b 45 03 01 05 00 00 00 00
    mov     %esp, %ebp
3  1: 89 e5 0c 8b 45 03 01 05 00 00 00 00
    mov     0xc(%ebp), %eax
4  3: 8b 45 03 01 05 00 00 00 00 00 00 00
    add     0x8(%ebp), %eax
5  6: 03 45 08
    add     %eax, 0x0
6  9: 01 05 00 00 00 00 00
    pop     %ebp
7  f: 5d
    ret
8  10: c3

```

Processor State

- The program counter (PC , %eip in IA32) indicates the address in memory of the next instruction to be executed.
 - The integer register file contains eight named locations storing 32-bit values. These registers can hold:
 - Addresses (corresponding to C pointers) or integer data.
 - Critical parts of the program state
 - Temporary data, such as the local variables of a procedure, and the value to be returned by a function.
-

Processor State

- The condition code registers hold status information about the most recently executed arithmetic or logical instruction.
 - These are used to implement conditional changes in the control or data flow, such as is required to implement *if* and *while* statements.
 - A set of floating-point registers store floating-point data.
-

Data Formats

- Intel uses the term “word” to refer to a 16-bit data type.
- Hence, 32- bit quantities as “double words” and 64-bit quantities as “quad words.”

C declaration	Intel data type	Assembly code suffix	Size (bytes)
char	Byte	b	1
short	Word	w	2
int	Double word	l	4
long int	Double word	l	4
long long int	—	—	4
char *	Double word	l	4
float	Single precision	s	4
double	Double precision	l	8
long double	Extended precision	t	10/12

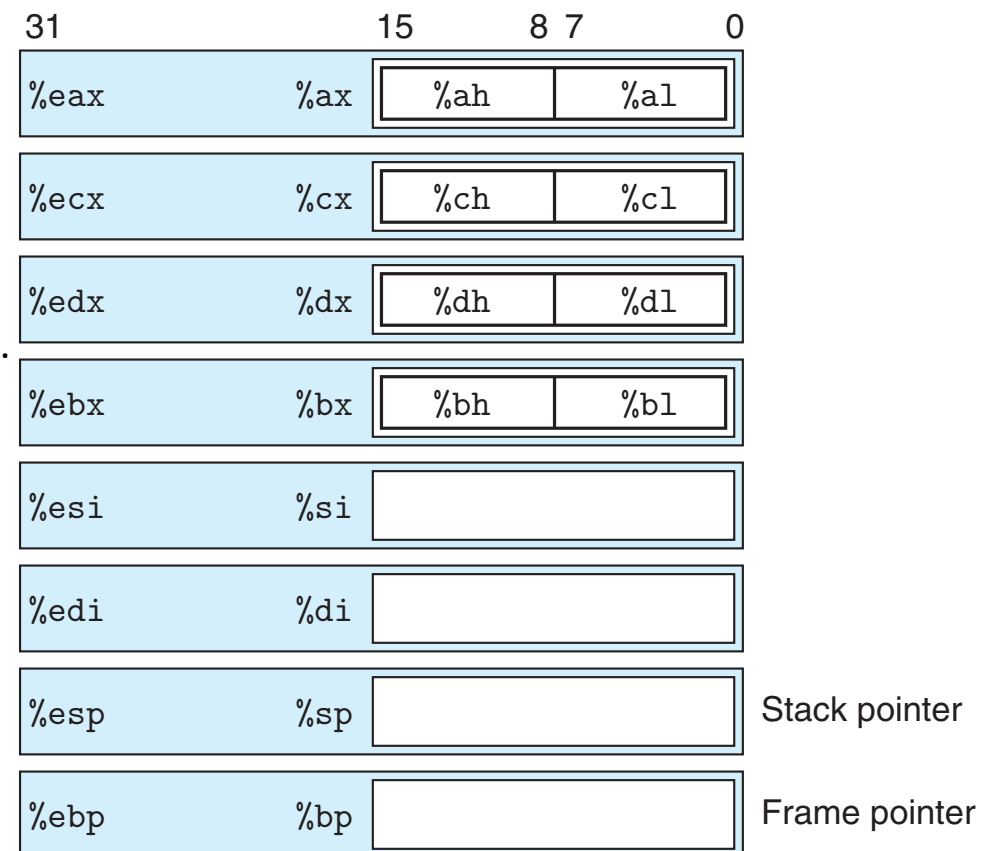
Sizes of C data types in IA32.

Data Formats

- IA32 does not support long long data type (recently introduced in C) in hardware. Instead, the compiler must generate sequences of instructions that operate on 32 bits data.
 - Most assembly-code instructions generated by gcc have a single-character suffix denoting the size of the operand e.g., movb (move byte), movw (move word), and movl (move double word).
 - Note that the assembly code uses the suffix 'l' to denote both a 4-byte integer as well as an 8-byte double-precision floating-point number.
 - This causes no ambiguity, since floating point involves an entirely different set of instructions and registers.
-

Accessing Information: IA32 Integer Registers

- 8 registers used to store integer/pointers
 - First six registers can be considered general-purpose registers, except some instructions that uses specific registers as source and/or destination.
 - Last two registers, contain pointers to important places in the program stack
 - The low-order 2 bytes of the first four registers can be independently read or written by the byte operation instructions.
 - Similarly, the low-order 16 bits of each register can be read or written by word operation instructions.



Accessing Information: x86-64 Integer Registers

- 16 registers to store integer/pointers
 - First six registers can be considered general-purpose registers, except some instructions that uses specific registers as source and/or destination.
 - Last two registers, contain pointers to important places in the program stack
 - The low-order 2 bytes of the first four registers can be independently read or written by the byte operation instructions.
 - Similarly, the low-order 16 bits of each register can be read or written by word operation instructions.

63	31	15	8	7	0	
%rax	%eax	%ax	%ah	%al		Return value
%rbx	%ebx	%bx	%bh	%bl		Callee saved
%rcx	%ecx	%cx	%ch	%cl		4th argument
%rdx	%edx	%dx	%dh	%dl		3rd argument
%rsi	%esi	%si		%sil		2nd argument
%rdi	%edi	%di		%dil		1st argument
%rbp	%ebp	%bp		%bpl		Callee saved
%rsp	%esp	%sp		%spl		Stack pointer
%r8	%r8d	%r8w		%r8b		5th argument
%r9	%r9d	%r9w		%r9b		6th argument
%r10	%r10d	%r10w		%r10b		Callee saved
%r11	%r11d	%r11w		%r11b		Used for linking
%r12	%r12d	%r12w		%r12b		Unused for C
%r13	%r13d	%r13w		%r13b		Callee saved
%r14	%r14d	%r14w		%r14b		Callee saved
%r15	%r15d	%r15w		%r15b		Callee saved

Operand Specifiers

- Most instructions have one or more operands, used as source and destination references

- Three types of operand exists
 - Constant (Immediate)
 - Register
 - Memory

Type	Form	Operand value	Name
Immediate	$\$Imm$	Imm	Immediate
Register	E_a	$R[E_a]$	Register
Memory	Imm	$M[Imm]$	Absolute
Memory	(E_a)	$M[R[E_a]]$	Indirect
Memory	$Imm(E_b)$	$M[Imm + R[E_b]]$	Base + displacement
Memory	(E_b, E_i)	$M[R[E_b] + R[E_i]]$	Indexed
Memory	$Imm(E_b, E_i)$	$M[Imm + R[E_b] + R[E_i]]$	Indexed
Memory	$(, E_i, s)$	$M[R[E_i] \cdot s]$	Scaled indexed
Memory	$Imm(, E_i, s)$	$M[Imm + R[E_i] \cdot s]$	Scaled indexed
Memory	(E_b, E_i, s)	$M[R[E_b] + R[E_i] \cdot s]$	Scaled indexed
Memory	$Imm(E_b, E_i, s)$	$M[Imm + R[E_b] + R[E_i] \cdot s]$	Scaled indexed

Operand Specifiers

Address	Value	Register	Value
0x100	0xFF	%eax	0x100
0x104	0xAB	%ecx	0x1
0x108	0x13	%edx	0x3
0x10C	0x11		

Operand	Value
%eax	_____
0x104	_____
\$0x108	_____
(%eax)	_____
4(%eax)	_____
9(%eax,%edx)	_____
260(%ecx,%edx)	_____
0xFC(,%ecx,4)	_____
(%eax,%edx,4)	_____

Data Movement Instructions

- Generally used for copy data from one location to another.
 - IA32/x86-64 imposes the restriction that a move instruction cannot have both operands refer to memory locations.
 - Copying a value from one memory location to another requires two instructions
 - First to load the source value into a register,
 - Second to write this register value to the destination.
-

Data Movement Instructions

Instruction		Effect	Description
MOV	S, D	$D \leftarrow S$	Move
movb		Move byte	
movw		Move word	
movl		Move double word	
MOVS	S, D	$D \leftarrow \text{SignExtend}(S)$	Move with sign extension

Data Movement Instructions

1	<code>movl \$0x4050,%eax</code>	<i>Immediate--Register, 4 bytes</i>
2	<code>movw %bp,%sp</code>	<i>Register--Register, 2 bytes</i>
3	<code>movb (%edi,%ecx),%ah</code>	<i>Memory--Register, 1 byte</i>
4	<code>movb \$-17, (%esp)</code>	<i>Immediate--Memory, 1 byte</i>
5	<code>movl %eax,-12(%ebp)</code>	<i>Register--Memory, 4 bytes</i>

Data Movement Instructions

Instruction		Effect	Description
movsbw		Move sign-extended byte to word	
movsbl		Move sign-extended byte to double word	
movswl		Move sign-extended word to double word	
MOVZ	S, D	$D \leftarrow \text{ZeroExtend}(S)$	Move with zero extension
movzbw		Move zero-extended byte to word	
movzbl		Move zero-extended byte to double word	
movzwl		Move zero-extended word to double word	

- In case of x86-64, we have additional set of instruction for quad word too, i.e., movsbq, movswq and movslq.
- However, for unsigned copying, only movzbq, movzwq exist as movl already set the high order 4 bytes to zero while copying 4 byte to 8 byte register.

Data Movement Instructions

Assume initially that %dh = CD, %eax = 98765432

`movb %dh,%al` *%eax = 987654CD*

`movsbl %dh,%eax` *%eax = FFFFFFFCD*

`movzbl %dh,%eax` *%eax = 000000CD*

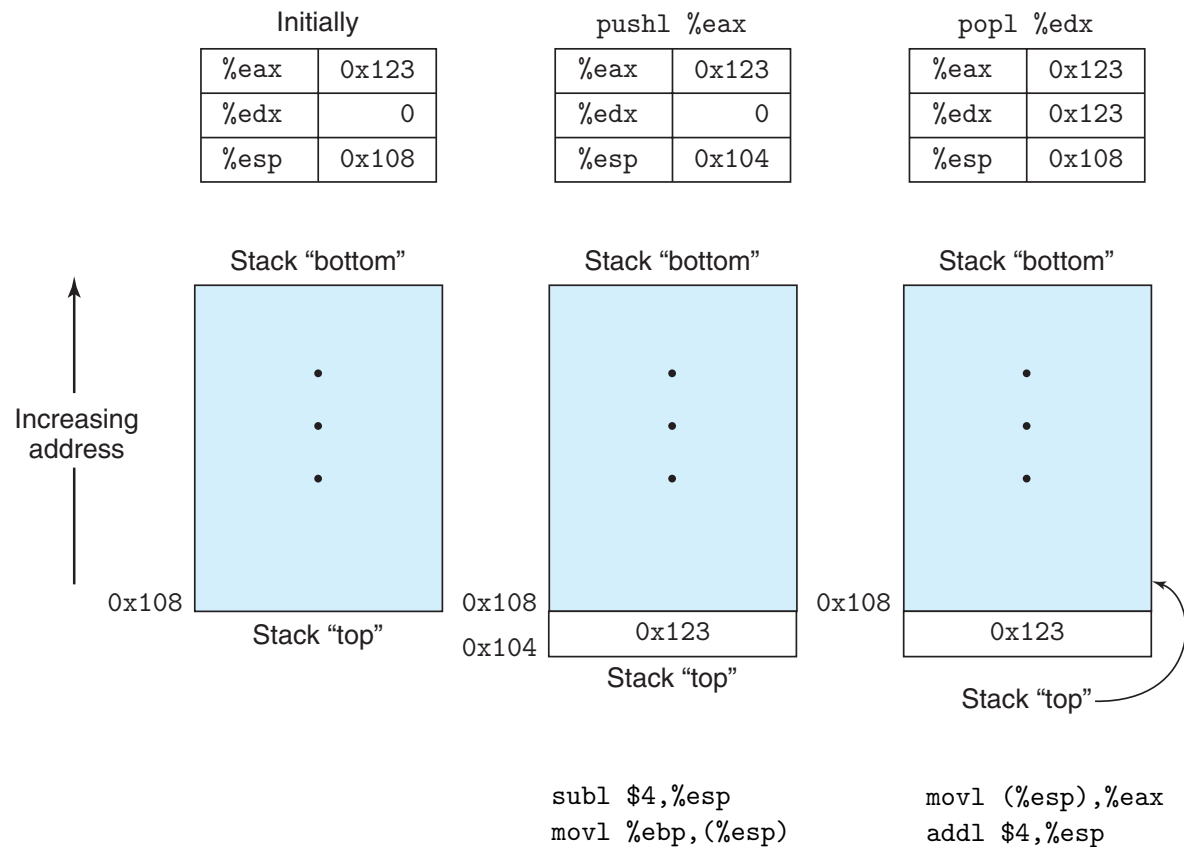
IA 32 Example

Data Movement Instructions

Instruction		Effect	Description
pushl	S	$R[\%esp] \leftarrow R[\%esp] - 4;$ $M[R[\%esp]] \leftarrow S$	Push double word
popl	D	$D \leftarrow M[R[\%esp]];$ $R[\%esp] \leftarrow R[\%esp] + 4$	Pop double word

Data Movement on Stack: Push & Pop

IA 32 Example



Data Movement Instructions Example

(a) C code

```
1  int exchange(int *xp, int y)
2  {
3      int x = *xp;
4
5      *xp = y;
6      return x;
7  }
```

(b) Assembly code

```
                                xp at %ebp+8, y at %ebp+12
1      movl    8(%ebp), %edx    Get xp
                                By copying to %eax below, x becomes the return value
2      movl    (%edx), %eax     Get x at xp
3      movl    12(%ebp), %ecx   Get y
4      movl    %ecx, (%edx)     Store y at xp
```
